

# l3p0fkite

February 7, 2025

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, classification_report, \
    ↪confusion_matrix
```

```
[3]: from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.naive_bayes import GaussianNB
```

```
[4]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

## 1 Loading the Dataset

```
[5]: dataset_path = "/content/sr_wq_rs_join.csv"
df = pd.read_csv(dataset_path, encoding='latin1')
```

```
[6]: unnecessary_columns = ['system:index', 'SiteID', 'date_unity', 'path', 'row', \
    ↪'sat', '.geo', 'endtime', 'date', 'date_only', 'source', 'lat', 'long', \
    ↪'TZID', 'date_utc', 'clouds', 'time', 'landsat_id', 'timediff', 'pwater', \
    ↪'type', 'id']
df = df.drop(columns=unnecessary_columns, errors='ignore')
```

```
[7]: df.fillna(df.median(), inplace=True)
```

## 2 Feature Selection

```
[8]: features = ['chl_a', 'doc', 'secchi', 'tss', 'p_sand']
```

```
[9]: def classify_suitability(value):
      if value >= 0.7:
          return "Highly Preferred"
      elif 0.4 <= value < 0.7:
          return "Partially Preferred"
      else:
          return "Least Preferred"
```

```
[10]: df['water_suitability'] = (df['chl_a'] * 0.3 + df['secchi'] * 0.3 + df['doc'] * 0.2 + df['tss'] * 0.2) / 4 # Example formula
      df['agriculture_suitability'] = df['water_suitability'].
      apply(classify_suitability)
```

### 3 Label Encoding

```
[11]: label_encoder = LabelEncoder()
      df['suitability_encoded'] = label_encoder.
      fit_transform(df['agriculture_suitability'])
```

```
[12]: X = df[features]
      y = df['suitability_encoded']
```

### 4 Model Training

```
[13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      random_state=42)
```

```
[14]: scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)
```

### 5 Random Forest

```
[15]: rf_model = RandomForestClassifier(n_estimators=200, max_depth=10,
      random_state=42)
      rf_model.fit(X_train, y_train)
      rf_preds = rf_model.predict(X_test)
      print("Random Forest Accuracy:", accuracy_score(y_test, rf_preds))
```

Random Forest Accuracy: 0.999187982135607

## 6 XGBoost Model

```
[16]: xgb_model = XGBClassifier(n_estimators=300, learning_rate=0.05, max_depth=10)
xgb_model.fit(X_train, y_train)
xgb_preds = xgb_model.predict(X_test)
print("XGBoost Accuracy:", accuracy_score(y_test, xgb_preds))
```

XGBoost Accuracy: 0.9994448449294456

## 7 Naive Bayes Model

```
[17]: nb_model = GaussianNB()
nb_model.fit(X_train, y_train)
nb_preds = nb_model.predict(X_test)
print("Naive Bayes Accuracy:", accuracy_score(y_test, nb_preds))
```

Naive Bayes Accuracy: 0.9836270683669326

## 8 Neural Networks

```
[19]: nn_model = Sequential()
nn_model.add(Dense(128, activation='relu', input_shape=(X_train.shape[1],)))
nn_model.add(Dropout(0.3))
nn_model.add(Dense(64, activation='relu'))
nn_model.add(Dropout(0.3))
nn_model.add(Dense(32, activation='relu'))
nn_model.add(Dense(3, activation='softmax'))
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87:  
UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When  
using Sequential models, prefer using an `Input(shape)` object as the first  
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[21]: from tensorflow.keras.callbacks import EarlyStopping
```

## 9 Optimization

```
[22]: nn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',  
    ↪ metrics=['accuracy'])  
early_stopping = EarlyStopping(monitor='val_loss', patience=5,  
    ↪ restore_best_weights=True)  
nn_model.fit(X_train, y_train, epochs=30, batch_size=32,  
    ↪ validation_data=(X_test, y_test), callbacks=[early_stopping])
```

```

Epoch 1/30
15086/15086          56s 3ms/step
- accuracy: 0.9980 - loss: 0.0053 - val_accuracy: 0.9989 - val_loss: 0.0029
Epoch 2/30
15086/15086          77s 3ms/step
- accuracy: 0.9982 - loss: 0.0056 - val_accuracy: 0.9982 - val_loss: 0.0042
Epoch 3/30
15086/15086          80s 3ms/step
- accuracy: 0.9981 - loss: 0.0050 - val_accuracy: 0.9988 - val_loss: 0.0033
Epoch 4/30
15086/15086          84s 3ms/step
- accuracy: 0.9981 - loss: 0.0049 - val_accuracy: 0.9987 - val_loss: 0.0032
Epoch 5/30
15086/15086          51s 3ms/step
- accuracy: 0.9983 - loss: 0.0045 - val_accuracy: 0.9990 - val_loss: 0.0027
Epoch 6/30
15086/15086          82s 3ms/step
- accuracy: 0.9984 - loss: 0.0052 - val_accuracy: 0.9991 - val_loss: 0.0026
Epoch 7/30
15086/15086          51s 3ms/step
- accuracy: 0.9981 - loss: 0.0048 - val_accuracy: 0.9990 - val_loss: 0.0035
Epoch 8/30
15086/15086          47s 3ms/step
- accuracy: 0.9984 - loss: 0.0050 - val_accuracy: 0.9991 - val_loss: 0.0028
Epoch 9/30
15086/15086          52s 3ms/step
- accuracy: 0.9984 - loss: 0.0042 - val_accuracy: 0.9987 - val_loss: 0.0030
Epoch 10/30
15086/15086          52s 3ms/step
- accuracy: 0.9983 - loss: 0.0046 - val_accuracy: 0.9985 - val_loss: 0.0037
Epoch 11/30
15086/15086          81s 3ms/step
- accuracy: 0.9984 - loss: 0.0047 - val_accuracy: 0.9982 - val_loss: 0.0041

```

[22]: <keras.src.callbacks.history.History at 0x7d97994c5190>

## 10 Evaluation of the Model

```

[23]: nn_preds = np.argmax(nn_model.predict(X_test), axis=1)
      print("Neural Network Accuracy:", accuracy_score(y_test, nn_preds))

```

```

3772/3772          7s 2ms/step
Neural Network Accuracy: 0.9991382667561544

```

## 11 Classification Report

```
[24]: print("\nRandom Forest Classification Report:\n", classification_report(y_test, rf_preds))
      print("\nXGBoost Classification Report:\n", classification_report(y_test, xgb_preds))
      print("\nNaive Bayes Classification Report:\n", classification_report(y_test, nb_preds))
      print("\nNeural Network Classification Report:\n", classification_report(y_test, nn_preds))
```

Random Forest Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	119601
1	0.88	0.35	0.50	20
2	0.96	0.94	0.95	1066
accuracy			1.00	120687
macro avg	0.95	0.76	0.82	120687
weighted avg	1.00	1.00	1.00	120687

XGBoost Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	119601
1	0.92	0.60	0.73	20
2	0.97	0.97	0.97	1066
accuracy			1.00	120687
macro avg	0.96	0.86	0.90	120687
weighted avg	1.00	1.00	1.00	120687

Naive Bayes Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	119601
1	0.04	1.00	0.08	20
2	0.29	0.59	0.39	1066
accuracy			0.98	120687
macro avg	0.45	0.86	0.49	120687
weighted avg	0.99	0.98	0.99	120687

# Neural Network Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	119601
1	0.00	0.00	0.00	20
2	0.96	0.94	0.95	1066
accuracy			1.00	120687
macro avg	0.65	0.65	0.65	120687
weighted avg	1.00	1.00	1.00	120687

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.11/dist-packages/sklearn/metrics/\_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))